Atomic adds under contention: Julia/Atomix vs Kokkos

RTX 4070 · N = 2^24 updates scattered onto M targets · median time per kernel launch (log–log, lower is better) · contention increases to the right

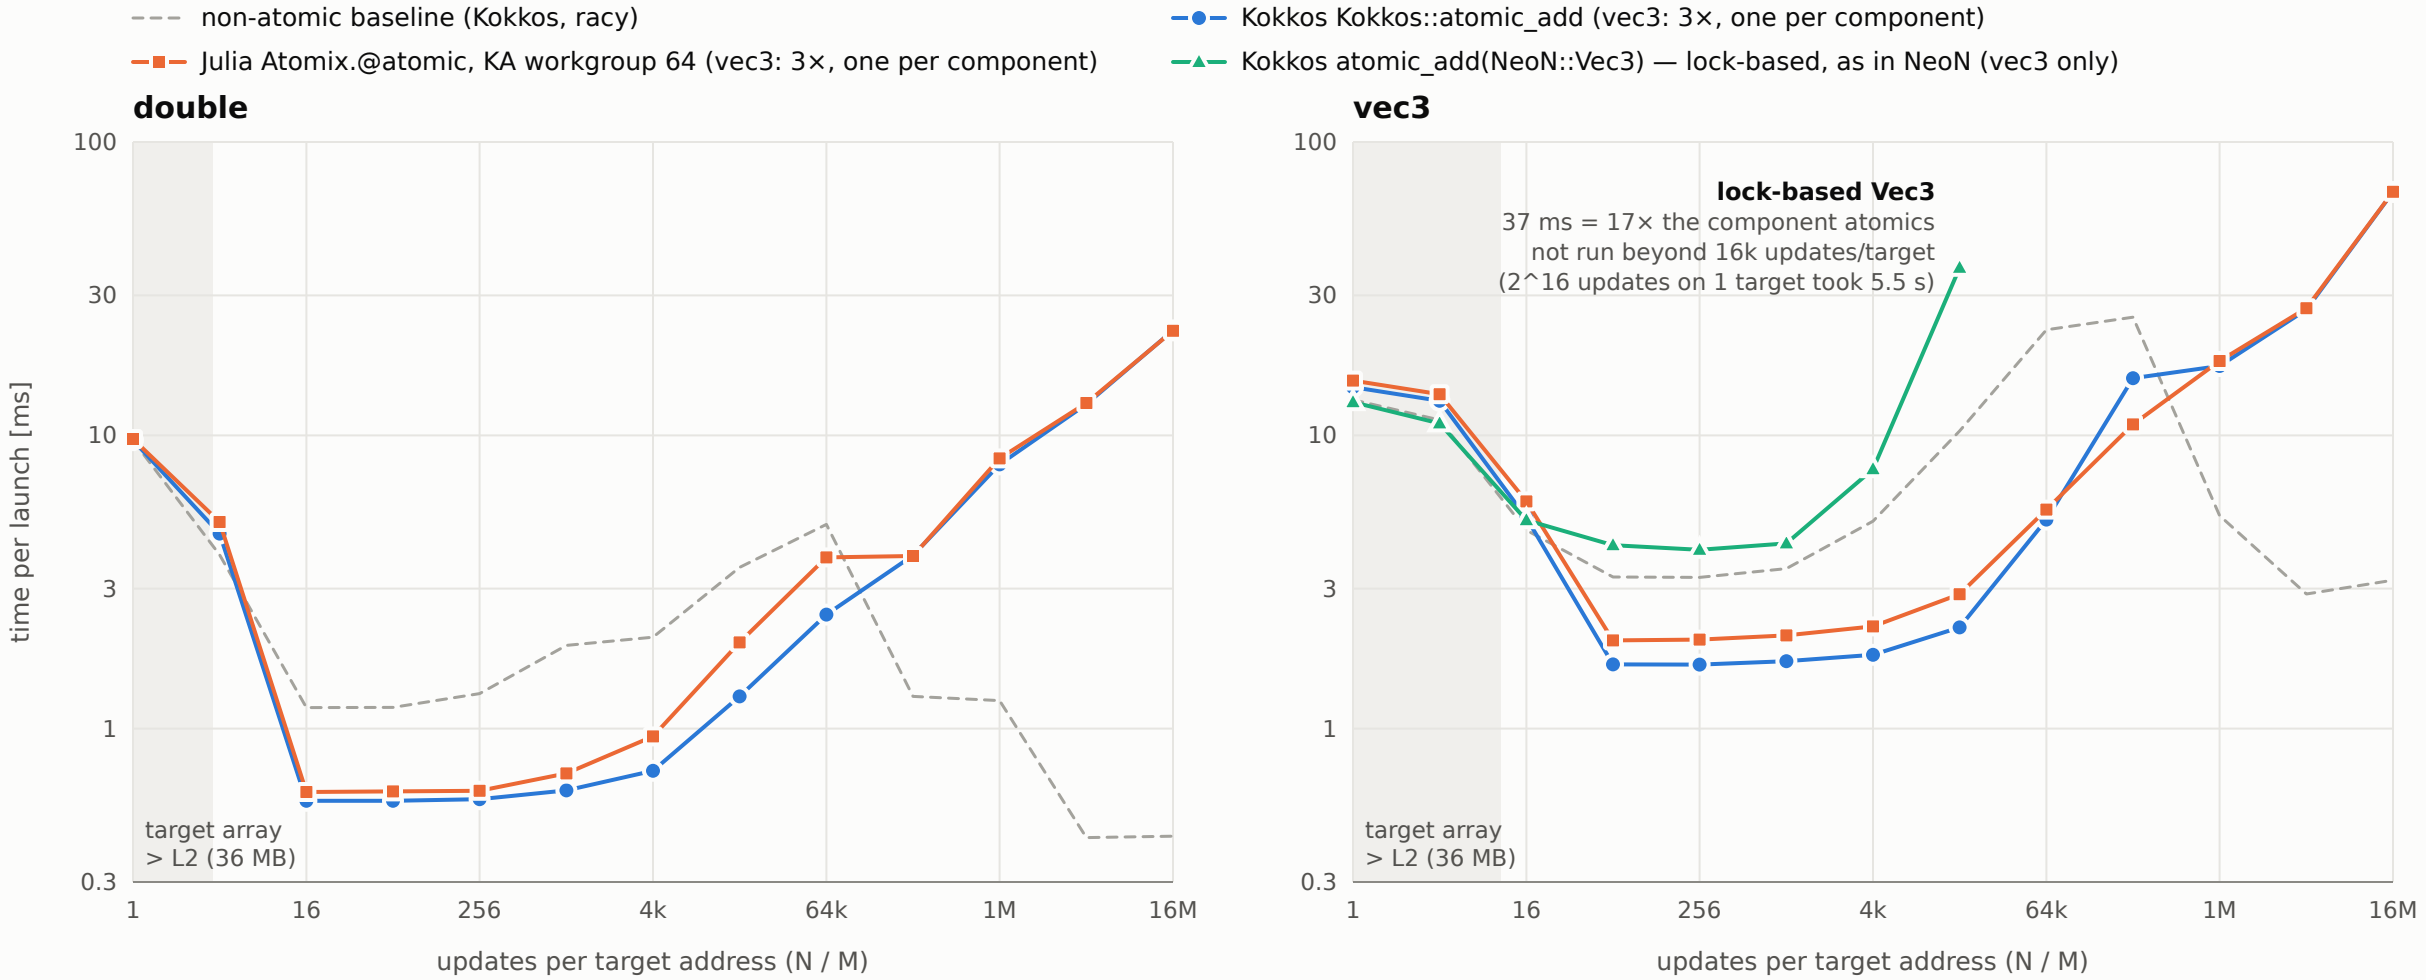

Julia with workgroup 256 is within noise of 64 (see summary.md). Vec3s are interleaved xyz doubles on both sides; same index pattern (splitmix64 % M) on both sides.